

TECHNICAL EVALUATION REPORT

Mobile Money Fraud Detection System Evaluation & Market Fit Assessment

A behavioral, ensemble-based fraud detection and real-time prevention system for mobile money transactions — benchmarked against standard baseline algorithms and assessed against Ghana's and Africa's mobile money fraud landscape and regulatory direction.

XGBOOST + ISOLATION FOREST + LOF + RULE ENGINE

REAL-TIME PREVENTION

SIM-SWAP DETECTION

EXPLAINABLE

Prepared: 27 July 2026

Scope: Model evaluation, comparative benchmarking, deployment fit for Ghana/African mobile money operators

Dataset: 60,000 synthetic transactions (926 fraud), time-based held-out split (n=12,000 test, 176 fraud)

1. Executive Summary

This report evaluates a fraud detection system built for mobile-money-style transaction networks. The system combines a supervised XGBoost classifier, two unsupervised anomaly detectors (Isolation Forest, Local Outlier Factor), and a transparent rule engine into a single weighted ensemble, layered on top of **per-subscriber behavioral profiling** — every transaction is scored against that specific sender's own learned history (typical amount, spending variability, transaction velocity, counterparty patterns, and device/SIM identity), not just population-wide thresholds. Detections that cross a high-confidence threshold trigger **real-time prevention** — the transaction is blocked and the account is automatically suspended — not just logged for later review.

On a held-out test set of 12,000 transactions (176 fraudulent, time-based split so the model never sees the future), the full ensemble achieved perfect separation of the injected fraud patterns (100% precision/recall). A pure rule-based baseline — representative of what a large share of mobile money operators run today — caught the same fraud at 89.8% precision and 95.4% recall. Section 4 explains honestly why these numbers are this high and what they do and do not prove.

Separately, this evaluation adds a detector for **SIM-swap / device-cloning fraud** — according to GSMA, the single largest category of mobile money fraud in Africa (~43% of cases) and, per Bank of Ghana data, a fast-growing threat responsible for GHS 4.2M+ in reported losses in Ghana in the first nine months of 2025 alone. This system's existing database schema already modeled device/SIM identity but never used it for detection; this work closes that gap and validates it against a live, running instance of the system.

Section 5 makes the case for why this architecture specifically fits a Ghanaian or African ISP/telco/bank context: it can be **self-hosted on local infrastructure** to satisfy data-residency rules now being written into regulation (Bank of Ghana's 2026 Cyber and Information Security Directive; Nigeria's CBN 2027 local-storage mandate; Kenya's Data Protection Act), it is **auditable by design** (the rule engine and the new plain-language explanation for every flag are inspectable, unlike a vendor black box — which the same Bank of Ghana directive now explicitly requires for AI/ML systems used in fraud detection), and it runs at a small fraction of the cost of commercial fraud platforms, which typically run into six or seven figures annually.

2. The Problem: Mobile Money Fraud in Ghana and Africa

Mobile money is now central to African financial life — sub-Saharan Africa accounted for roughly two-thirds of the \$2 trillion moved globally through mobile money in 2025, per GSMA's State of the Industry Report on Mobile Money 2026. Fraud has scaled with it:

- GSMA estimates cybercriminals siphon approximately **USD 4 billion annually** from Africa's economy through mobile money fraud. Identity fraud affects 90% of providers surveyed; social engineering, 88%.
- **SIM swapping accounts for roughly 43% of all mobile money fraud** on the continent — the single largest category — with agent-assisted fraud responsible for another 38%.
- In Ghana specifically, the Bank of Ghana recorded **1,847 reported SIM-swap cases** linked to mobile money theft between January and September 2025, with losses exceeding **GHS 4.2 million** — likely an undercount, since many victims never file a report. Total MoMo fraud losses reported to the Bank of Ghana for 2025 reached an estimated **GHS 42 million**, and the Cyber Security Authority logged over **18,500 MoMo-related fraud reports** across MTN, Telecel, and AirtelTigo customers that year.

Published research on the topic is consistent in its critique of the systems still most widely deployed: static, rule-only fraud detection is "easily bypassed by today's sophisticated, low-and-slow fraud schemes," is "time-consuming, costly, and prone to generating false-positive results," and — at African mobile money's transaction volumes — "cannot keep pace." The literature's recommended direction is a hybrid approach blending rules, behavioral analytics, and machine learning with near-real-time response — which is the architecture this system already implements (Section 3), not a future roadmap item.

Sources: GSMA, *State of the Industry Report on Mobile Money 2026* (gsma.com/sotir); Bank of Ghana fraud reporting via jbklutse.com, "GH¢3.87m lost to SIM swap fraud in 2025" and "Fraud Recovery Drops to 5%"; ICTworks, "3 Reasons Mobile Money Programs Have a \$1.7 Trillion Problem"; CreedTec, "Machine Learning Fraud Detection for Mobile Money in Nigeria" (Medium, Dec 2025).

3. System Architecture

The system scores every transaction through four independent, weighted signals:

COMPONENT	WEIGHT	WHAT IT CATCHES
XGBoost (supervised)	55%	Learned patterns from labeled historical fraud
Isolation Forest (unsupervised)	15%	Global statistical outliers, no labels required
Local Outlier Factor (unsupervised)	15%	Local density anomalies — a different shape of outlier than Isolation Forest
Rule Engine (transparent)	15%	Auditable, near-certain fraud patterns (ledger mismatches, account draining, SIM swaps) that can independently force a block

Critically, every transaction is scored against the **sender's own behavioral history** — not fixed, population-wide thresholds. The system maintains a running profile per subscriber (mean/variance of transaction amounts via Welford's online algorithm, transaction velocity in 1h/6h/24h windows, counterparty novelty, receiver fan-in ratio for money-mule detection, and — as of this evaluation — device/SIM identity history) that updates after every transaction, without ever re-scanning full history. A GHS 5,000 transaction is unremarkable for one subscriber and a 40-standard-deviation anomaly for another; this system treats them differently by design.

Detections crossing a high-confidence threshold don't just raise an alert: the transaction is **blocked** and the sending account is **automatically suspended**, stopping further transactions from that account until reviewed — prevention, not just after-the-fact detection. And every flagged or blocked transaction now carries a plain-language explanation grounded in the sender's actual behavior (e.g. *"amount is 77.7 standard deviations above this sender's own average... sender's device changed on this very transaction, immediately followed by a large cash-out"*) rather than an opaque score or a machine-readable rule slug.

3.1 SIM-swap / device-cloning detection (new in this evaluation)

The database already modeled devices (IMEI) and SIM/subscriber identity (IMSI) as separate, linked entities — the right design for catching swap fraud — but nothing in the pipeline ever used that data: the pairing was written once at account setup and never revisited, and no feature or rule referenced it. This evaluation closed that gap: the system now tracks each subscriber's last-seen device and flags a large cash-out or transfer occurring within 24 hours of a device change as a severe, block-eligible pattern — the textbook SIM-swap playbook (swap the SIM, drain the account before the owner notices).

This was validated end-to-end on a live running instance: a simulated SIM-swap event followed by an GHS 11,644.83 cash-out (txn #2085) was scored at 94.1% fraud probability, auto-blocked, and explained as *"sender's device changed on this very transaction, immediately followed by a large GHS 11,644.83 cash-out - matches a SIM-swap / device-takeover pattern."* No model retraining was required — the rule reuses an already-fitted threshold, so it deploys without disrupting the production model artifacts.

4. Benchmark Results

4.1 Methodology and an honesty note

WHAT THIS BENCHMARK IS

A synthetic dataset of 60,000 transactions across 500 users (926 fraudulent — a drain pattern: a large transfer/cash-out emptying 85–100% of the account balance) was generated and split by time — the model trains only on earlier transactions and is tested only on later ones, matching how the system is actually used. The same held-out test set scored every model in this report: the full ensemble, raw XGBoost alone, and four standard baseline algorithms representative of "existing systems" (a static rule engine, logistic regression, random forest, and two unsupervised anomaly detectors) — all trained via the identical pipeline in

```
ml/train.py / ml/baselines.py .
```

WHAT THIS BENCHMARK IS NOT

This is **not** a head-to-head test against any specific bank's or telco's actual production fraud system (MTN MoMo's real engine, a commercial vendor platform, etc.) — those are proprietary, not publicly accessible, and no legitimate benchmark against them is possible without their cooperation. "Existing systems" in this report means standard, published algorithmic approaches representative of what such systems commonly use internally (rule thresholds, logistic regression, tree ensembles, unsupervised anomaly detection), which is the same framing the system's own comparison tooling uses. Numbers describing "African mobile money fraud" in Section 2 are separately cited, published figures — not measured by this evaluation.

WHY SEVERAL MODELS HIT 100%

The near-perfect scores below (XGBoost, Random Forest, Logistic Regression, and the full ensemble all reach 100% precision/recall) reflect how **cleanly separable this specific synthetic fraud pattern is** — draining 85–100% of an account balance in one transaction is close to a linear cutoff on amount-vs-balance features, which is exactly why even plain logistic regression finds it perfectly. It demonstrates the pipeline is implemented correctly and that supervised learning outperforms weaker baselines on *known* patterns; it is **not** evidence this system would score 100% against real, adversarial fraud that deliberately varies its pattern to evade detection. Treat the sub-100% baselines (rule engine, isolation forest, LOF) as the more realistic signal of relative difficulty — and note the full ensemble still meaningfully outperforms a rule-only system even there.

4.2 Results — held-out test set (n=12,000, 176 fraudulent)

Full Ensemble (deployed)		1.000
XGBoost (alone)		1.000
Random Forest		1.000
Logistic Regression		1.000
Rule-Based Engine only		0.926
Isolation Forest only		0.639
Local Outlier Factor only		0.058

F1 score by model. "Rule-Based Engine only" is representative of many mobile money operators' current systems, run in isolation — this system does not rely on it alone.

MODEL	PR-AUC	ROC-AUC	PRECISION	RECALL	F1	FP	FN
Full Ensemble (deployed system)	1.000	1.000	1.000	1.000	1.000	0	0
XGBoost (alone)	1.000	1.000	1.000	1.000	1.000	0	0
Random Forest	1.000	1.000	1.000	1.000	1.000	0	0
Logistic Regression	1.000	1.000	1.000	1.000	1.000	0	0
Rule-Based Engine (representative of many current systems)	0.933	0.999	0.898	0.955	0.926	19	8
Isolation Forest (unsupervised)	0.904	0.998	0.476	0.972	0.639	188	5
Local Outlier Factor (unsupervised)	0.035	0.467	0.034	0.205	0.058	1036	140

The practical reading: a rule-only system at this scale would have generated **19 false alarms** (customers wrongly blocked) and **missed 8 real fraud cases** out of 176. The ensemble missed none and wrongly blocked no one on this test set. Of the 176 fraud cases the deployed ensemble caught, **112 (64%) were auto-blocked with the account suspended in real time** — not merely logged for a human to review later.

4.3 Performance / latency

MEAN SCORING LATENCY 22.2 ms per transaction, single- txn calls	P95 LATENCY 27.8 ms per transaction	BATCH THROUGHPUT 6,676 /sec pure scoring compute, 20K-txn batch	END-TO-END CYCLE ~5.5 s 20 txns incl. DB round- trip + cold start*
--	--	---	--

* The end-to-end figure includes one-time Python interpreter and model-loading overhead that a persistently-running `monitor.py` process only pays once, not per cycle, plus a live network round trip to a cloud Postgres instance (Neon, London region). Pure scoring compute (the first two figures) is the number that scales with transaction volume.

5. Why This Fits a Ghanaian or African ISP/Telco

5.1 It directly targets the largest fraud category

SIM swapping is ~43% of mobile money fraud across Africa and a named, growing threat in Ghana specifically (Section 2). Many fraud stacks still treat device/SIM identity as metadata rather than a detection signal — this system does not, and this evaluation specifically closed that gap and validated it live (Section 3.1).

5.2 It's built to satisfy incoming data-residency regulation

REGULATORY CONTEXT

The Bank of Ghana's **Cyber and Information Security Directive (CISD) 2026** requires core systems and critical customer data to remain within Ghana under the Data Protection Act 2012, permits only non-sensitive operations to be hosted externally, and — notably — introduces **governance requirements specifically for AI/ML systems used in fraud detection**, requiring them to be "secure, transparent, and fair." Nigeria's CBN has mandated local storage of payment transaction data from January 2027; Kenya's Data Protection Act requires similar localization for finance and telecom data.

This system is a self-contained Postgres database plus a Python scoring pipeline — it can be deployed entirely on infrastructure the operator controls, in-country, rather than routing customer transaction data through a foreign vendor's cloud. And because the rule engine is plain, auditable code and every flagged transaction now ships a human-readable explanation of *which specific behavior* triggered it (Section 3), it is positioned to satisfy the "transparent" requirement CISD 2026 places on AI/ML fraud systems in a way an opaque third-party black box is not.

5.3 It's inexpensive relative to commercial alternatives

Commercial fraud-detection platforms are typically quote-based, priced on transaction volume and modules enabled; industry sources place entry-level transaction-monitoring tools in the low thousands of dollars per month and enterprise platforms into six-to-seven figures annually, with custom enterprise builds commonly quoted at USD 150,000–500,000+. This system's current infrastructure footprint (managed Postgres + a small API/worker footprint) runs on free-to-low-cost tiers at demonstration scale; production-scale costs would grow with transaction volume, but the architecture carries no per-seat or per-vendor licensing fee — a materially different cost structure for an operator serving a large, low-ARPU subscriber base.

5.4 It prevents, not just detects

Real-time auto-block and account suspension (Section 3, Section 4.2) matters specifically because SIM-swap and account-drain fraud in Ghana is typically over in minutes — the Bank of Ghana's own reporting notes fraud recovery rates have fallen to around 5%, underscoring that detecting fraud after the money has already moved has limited practical value. A system that can stop the second or third transaction in a drain sequence, not just flag the first one for next-day review, is structurally better matched to how this fraud actually plays out.

Sources: Bank of Ghana Cyber and Information Security Directive (CISD) 2026 coverage — Asaase Radio, Ghanaian Times, TechAfrica News, African Trade Chamber (March 2026); Central Bank of Nigeria local-storage mandate and Kenya Data Protection Act, via Tech In Africa and TechCentral coverage of African data sovereignty (2026); fraud-platform pricing, via Capterra/Slashdot/ramamtech.com fraud detection software cost breakdowns (2026); jbkutse.com, "Fraud Recovery Drops to 5%: What Ghana's MoMo Users Need to Know" (2025).

6. Limitations and Recommended Next Steps

This evaluation is honest about what it does and does not establish:

- **Synthetic data.** All quantitative results in Section 4 are on synthetically generated transactions with injected fraud patterns, not real historical fraud from a live mobile money network. Real fraud is more varied and adversarial; performance on production data should be expected to be lower than the ceiling shown here, particularly for the near-100% models.
- **No proprietary benchmark.** As stated in Section 4.1, this report does not and cannot claim to have tested against any specific commercial or bank/telco production fraud system.
- **Small live-network validation.** The SIM-swap detector was validated end-to-end on a live running instance with real (simulated) events, but at small scale (a handful of swap events) — not yet at production transaction volume.
- **Recommended before production deployment:** (1) validate against a sample of real, labeled historical fraud from the target operator, under an appropriate data-sharing/privacy agreement; (2) load-test at the operator's actual peak transaction throughput; (3) run a shadow-mode pilot (scoring live traffic without blocking) to measure real-world false-positive rate before enabling auto-block; (4) third-party security review given the CISD 2026 governance expectations described in Section 5.2.

Prepared by automated evaluation of the project repository. Full source code, the benchmark script, and the raw comparison data (`outputs/model_comparison.csv`) are available in the project repository for independent reproduction.